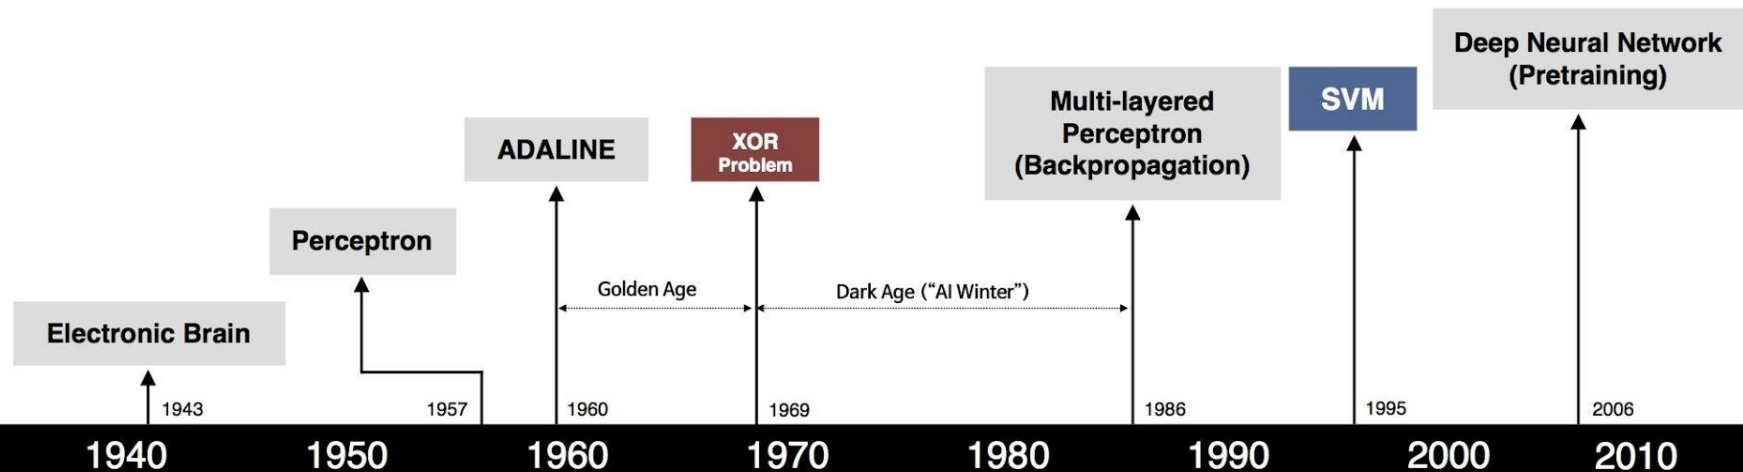


Нейронные сети


Лекция 7

История

1. 1940-ые – Идея перцептрона и искусственного мозга
2. 1950-ые – Мат. модель перцептрона
3. 1950-ые – Adaline модель
4. 1969 – Озвучена XOR проблема
5. 1986 – Многослойная модель, Backpropagation method
6. 1988 – Сверточные нейронные сети
7. 1990-ые – Рекуррентные нейронные сети
8. 2006 – Глубокие нейросети
9. 2009 – GPU для нейросетей
10. 2010-ые Image Net эпопея
11. ...




S. McCulloch – W. Pitts



- Adjustable Weights
- Weights are not Learned



F. Rosenblatt




- Learnable Weights and Threshold



B. Widrow – M. Hoff




M. Minsky – S. Papert



- XOR Problem




D. Rumelhart – G. Hinton – R. Williams



- Solution to nonlinearly separable problems
- Big computation, local optima and overfitting




V. Vapnik – C. Cortes




- Limitations of learning prior knowledge
- Kernel function: Human Intervention



G. Hinton – S. Ruslan



- Hierarchical feature Learning



Задача MNIST

- 10 классов (0..9)
- 28x28 пикселя
- 60_000 тренировочная
- 10_000 тестовая



28x28
pixels



784 pixels

"neurons"




0

1

2


9



$$L_n = \sum_{j=1}^M x_{i,j} * w_{n,j} + b_n = w_n^T x_i$$



$$\textit{Softmax} (L_n) = P (Y = n | X = x_i)$$

$$\textit{Softmax} (L_n) = \frac{e_n^L}{||e^L||} = \frac{e_n^L}{\sum_{i=0}^N e_i^L}$$



X: 100 images,
one per line,
flattened

8
4
6
6
4
3



10 columns

$W_{0,0}$	$W_{0,1}$	$W_{0,2}$	$W_{0,3}$...	$W_{0,9}$
$W_{1,0}$	$W_{1,1}$	$W_{1,2}$	$W_{1,3}$...	$W_{1,9}$
$W_{2,0}$	$W_{2,1}$	$W_{2,2}$	$W_{2,3}$...	$W_{2,9}$
$W_{3,0}$	$W_{3,1}$	$W_{3,2}$	$W_{3,3}$...	$W_{3,9}$
$W_{4,0}$	$W_{4,1}$	$W_{4,2}$	$W_{4,3}$...	$W_{4,9}$
$W_{5,0}$	$W_{5,1}$	$W_{5,2}$	$W_{5,3}$...	$W_{5,9}$
$W_{6,0}$	$W_{6,1}$	$W_{6,2}$	$W_{6,3}$...	$W_{6,9}$
$W_{7,0}$	$W_{7,1}$	$W_{7,2}$	$W_{7,3}$...	$W_{7,9}$
$W_{8,0}$	$W_{8,1}$	$W_{8,2}$	$W_{8,3}$...	$W_{8,9}$
...
$W_{783,0}$	$W_{783,1}$	$W_{783,2}$	$W_{783,9}$

784 lines

$L_{0,0}$	$L_{0,1}$	$L_{0,2}$	$L_{0,3}$...	$L_{0,9}$
$L_{1,0}$	$L_{1,1}$	$L_{1,2}$	$L_{1,3}$...	$L_{1,9}$
$L_{2,0}$	$L_{2,1}$	$L_{2,2}$	$L_{2,3}$...	$L_{2,9}$
$L_{3,0}$	$L_{3,1}$	$L_{3,2}$	$L_{3,3}$...	$L_{3,9}$
$L_{4,0}$	$L_{4,1}$	$L_{4,2}$	$L_{4,3}$...	$L_{4,9}$
...
$L_{99,0}$	$L_{99,1}$	$L_{99,2}$	$L_{99,9}$

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} = \text{softmax} \left(\begin{bmatrix} W_{1,1} & W_{1,2} & W_{1,3} \\ W_{2,1} & W_{2,2} & W_{2,3} \\ W_{3,1} & W_{3,2} & W_{3,3} \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix} \right)$$

Predictions

$Y[100, 10]$

Images

$X[100, 784]$

Weights

$W[784, 10]$

Biases

$b[10]$


$$Y = \text{softmax}(X \cdot W + b)$$

applied line
by line

matrix multiply

broadcast
on all lines

tensor shapes in []




0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	0	0	1	0

$$\textit{Softmax}(L_n)$$

0,1	0,2	0,1	0,5	0,4	0,1	0,3	0,2	0,1	0,6
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----

На сколько хорошо?




0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	0	0	1	0

Softmax (L_n)

0,1	0,2	0,1	0,5	0,4	0,1	0,3	0,2	0,1	0,6
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----


$$Loss(X_i) = - \sum Y_i' \log(Y_i)$$



0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	0	0	1	0
0,1	0,2	0,1	0,5	0,4	0,1	0,3	0,2	0,1	0,6

$$Loss(X_i) = -\log(0.1) = 2.30$$






0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	0	0	1	0
0,1	0,2	0,1	0,5	0,4	0,1	0,3	0,2	0,1	0,6

$$Loss(X_i) = -\log(0.7) = 0.36$$



Минимизация ошибки

Градиентный спуск



Оптимизаторы

- SGD
 - Adagrad
 - Adadelata
 - Adam
 - Adamax
 - Nesterov Adam
 - Rmsprop
 -
- https://www.tensorflow.org/api_guides/python/train#Optimizers
 - <https://keras.io/optimizers/>
 - <https://habrahabr.ru/post/318970/>

Однослойная нейросеть на TF

```
X = tf.placeholder(tf.float32, [None, 784], name='X')
W = tf.Variable(tf.zeros([784, 10]))
b = tf.Variable(tf.zeros([10]))
```

```
# Модель  $Y = \text{softmax}(X.W + b)$ 
```

```
Y = tf.nn.softmax(tf.matmul(
    tf.reshape(X, [-1, 784]), W) + b)
```

```
# Подстановка для корректных значений входных данных
```

```
Y_ = tf.placeholder(tf.float32, [None, 10])
```

```
# Оптимизируем функцию потерь методом градиентного спуска  
# 0.003 - это шаг градиента, гиперпараметр  
optimizer = tf.train.GradientDescentOptimizer(0.003)
```

```
# Минимизируем потери  
train_step = optimizer.minimize(cross_entropy)
```

```
# Функция потерь  $H = \text{Sum}(Y_ * \log(Y))$   
cross_entropy = -tf.reduce_sum(Y_ * tf.log(Y))
```

```
# доля корректных ответов  
is_correct = tf.equal(tf.argmax(Y,1), tf.argmax(Y_,1))  
# точность предсказания  
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))
```

инициализация сессии

```
init = tf.global_variables_initializer()
```

```
sess = tf.Session()
```

```
sess.run(init)
```

acc, entropy = [], [] # переменные для рисования графиков

```
for i in range(1000):
```

load batch of images and correct answers

```
batch_X, batch_Y = mnist.train.next_batch(100)
```

```
train_data={X: batch_X, Y_: batch_Y}
```

train

```
sess.run(train_step, feed_dict=train_data)
```


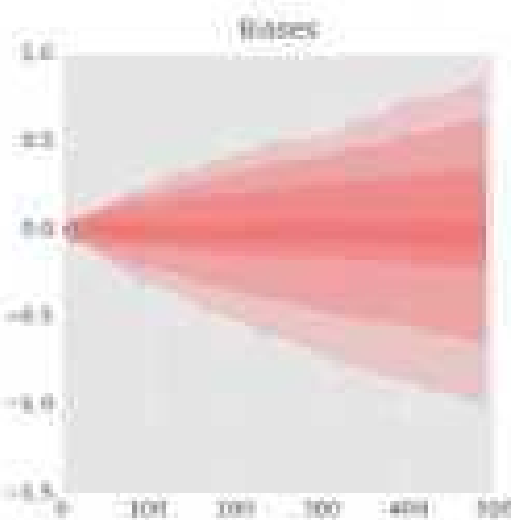
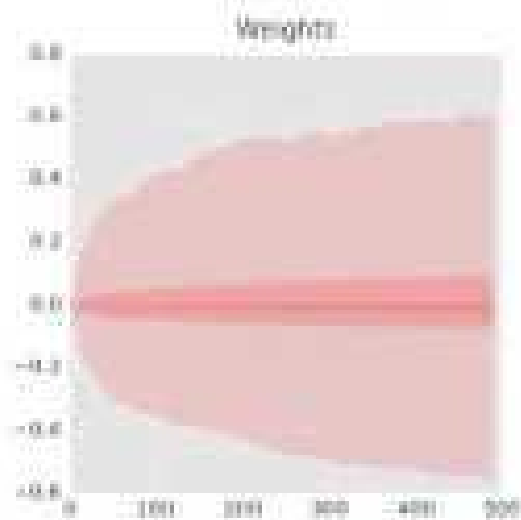
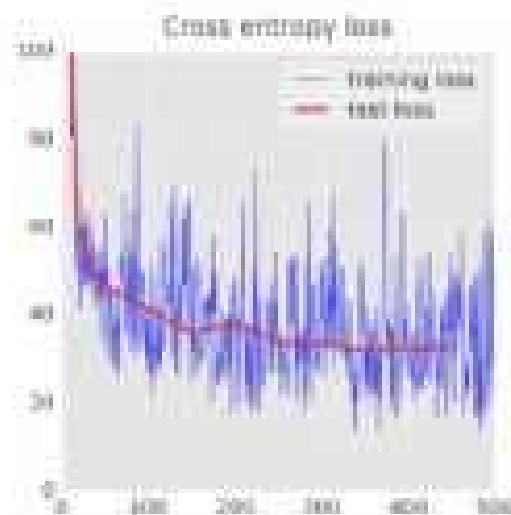
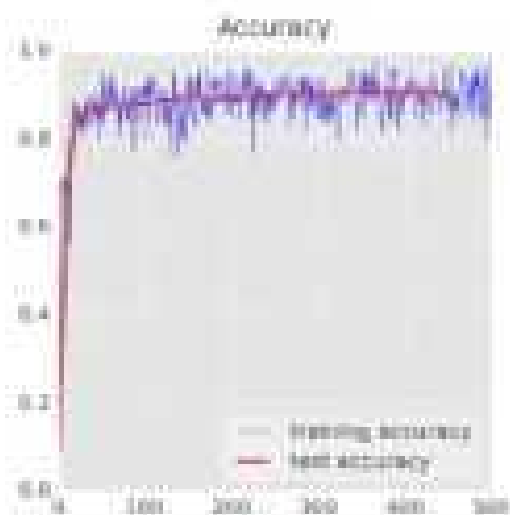
```
if i % 10 == 0:
```

```
    test_data={X: mnist.test.images, Y_: mnist.test.labels}
```


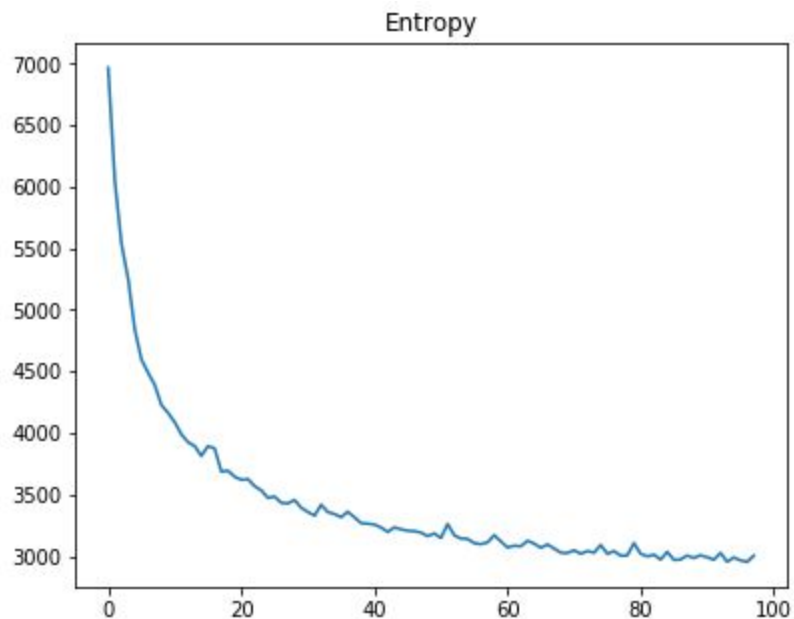
```
    a, c = sess.run([accuracy, cross_entropy], feed_dict=test_data)
```

```
    acc.append(a)
```

```
    entropy.append(c)
```



Результат однослойной сети (0.91)





A close-up shot from the movie Inception featuring Leonardo DiCaprio. He is wearing a dark suit and tie, looking slightly to his right with a serious expression. Another person's face is partially visible on the right side of the frame, looking towards DiCaprio. The background is blurred, showing what appears to be an office or meeting room setting.

WE NEED TO GO


DEEPER

Многослойная нейронная сеть





$$\textit{Sigmoid} = \frac{1}{1 + e^{-x}}$$



Объявляем веса

```
layer_sizes = [28*28, 200, 100, 60, 30, 10]

W1 = tf.Variable(tf.truncated_normal([layer_sizes[0], layer_sizes[1]], stddev=0.1))
B1 = tf.Variable(tf.zeros(layer_sizes[1])) # 200

W2 = tf.Variable(tf.truncated_normal([layer_sizes[1], layer_sizes[2]], stddev=0.1))
B2 = tf.Variable(tf.zeros(layer_sizes[2])) # 100

W3 = tf.Variable(tf.truncated_normal([layer_sizes[2], layer_sizes[3]], stddev=0.1))
B3 = tf.Variable(tf.zeros(layer_sizes[3])) # 60

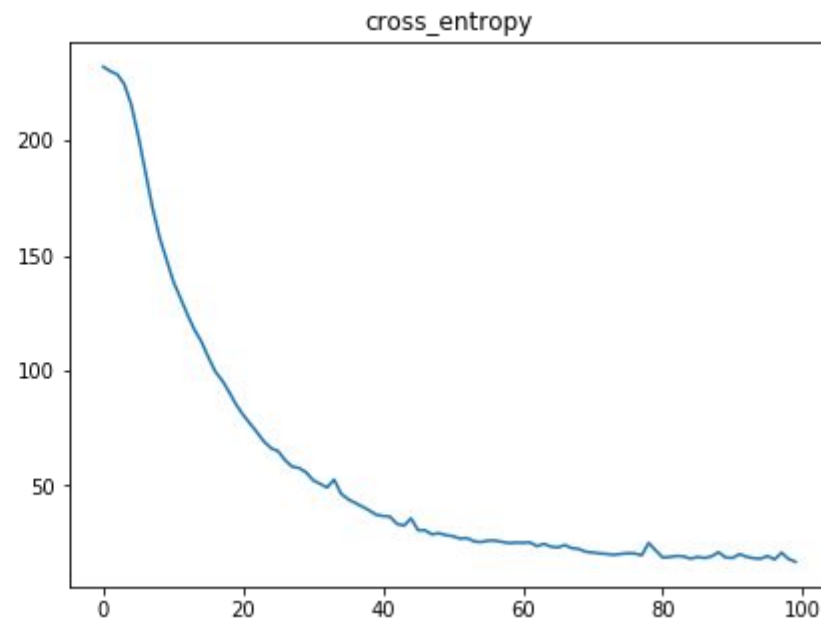

W4 = tf.Variable(tf.truncated_normal([layer_sizes[3], layer_sizes[4]], stddev=0.1))
B4 = tf.Variable(tf.zeros(layer_sizes[4])) # 30

W5 = tf.Variable(tf.truncated_normal([layer_sizes[4], layer_sizes[5]], stddev=0.1))
B5 = tf.Variable(tf.zeros(layer_sizes[5])) # 10
```

```
X = tf.placeholder(tf.float32, [None, 784])


Y1 = tf.nn.sigmoid(tf.matmul(X, W1) + B1)
Y2 = tf.nn.sigmoid(tf.matmul(Y1, W2) + B2)
Y3 = tf.nn.sigmoid(tf.matmul(Y2, W3) + B3)
Y4 = tf.nn.sigmoid(tf.matmul(Y3, W4) + B4)

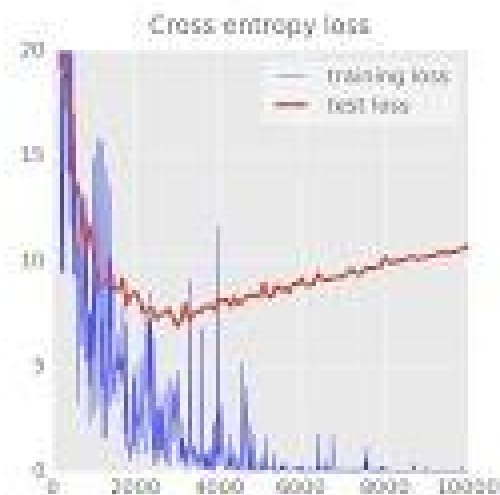
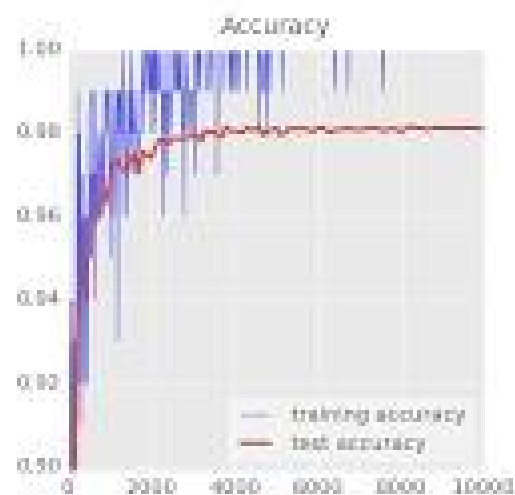
Ylogits = tf.matmul(Y4, W5) + B5
Y = tf.nn.softmax(Ylogits)
Y_ = tf.placeholder(tf.float32, [None, 10])
```



Медленное обучение
Нестабильность

Relu = Rectified Linear Unit







Training digits

```

5 0 0 2 4 6 7 8 3 9
4 0 1 7 5 6 9 6 0 6
0 6 1 0 0 6 0 6 3 9
8 1 7 4 6 3 6 9 5 9
0 5 1 6 3 2 7 5 7 3
4 7 0 8 1 5 5 2 7 1
6 4 4 2 1 3 6 7 3 4
0 7 8 0 7 0 1 1 7 4
3 7 7 0 3 8 0 1 8 9
4 4 7 9 5 0 8 1 6 1
  
```



Далее

- Learning rate decay
- Dropout for regularization
- Convolutional NN
- Deep conv NN

https://nbviewer.jupyter.org/github/mlomsk/ML_MNIST_Workshop_2017-09-28/blob/master/main.ipynb

Литература

- <http://cs231n.github.io/>
- <https://codelabs.developers.google.com/codelabs/cloud-tensorflow-mnist/>
- <http://neuralnetworksanddeeplearning.com/>
- http://sebastianraschka.com/Articles/2015_singlelayer_neurons.html
- <http://www.deeplearningbook.org/>
- <https://medium.com/@ageitgey/machine-learning-is-fun-part-2-a26a10b68df3>
- <https://keras.io/>
- <https://www.tensorflow.org/>